

Analysis of the conflict between omission and commission in low spatial resolution dichotomic thematic products: The Pareto Boundary

Luigi Boschetti^{a,*}, Stéphane P. Flasse^b, Pietro A. Brivio^c

^aGlobal Vegetation Monitoring Unit, European Commission, DG Joint Research Centre, Institute for Environment and Sustainability, via Fermi 1, I-21020 Ispra, VA, Italy

^bFlasse Consulting, 3 Sycamore Crescent, Maidstone, Kent ME16 0AG, UK

^cCNR-Istituto per il Rilevamento Elettromagnetico dell'Ambiente, via Bassini 15, I-20133 Milan, Italy

Received 8 July 2003; received in revised form 26 February 2004; accepted 28 February 2004

Abstract

During the last few years, the remote sensing community has been trying to address the need for global synthesis to support policy makers on issues such as deforestation or global climate change. Several global thematic products have been derived from large datasets of low-resolution remotely sensed data, the latter providing the best trade-off between spatial resolution, temporal resolution and cost. However, a standard procedure for the validation of such products has not been developed yet. This paper proposes a methodology, based on statistical indices derived from the widely used Error Matrix, to deal with the specific issue of the influence of the low spatial resolution of the dataset on the accuracy of the end-product, obtained with hard classification approaches. In order to analyse quantitatively the trade-off between omission and commission errors, we suggest the use of the 'Pareto Boundary', a method rooted in economics theory applied to decisions with multiple conflicting objectives. Starting from a high-resolution reference dataset, it is possible to determine the maximum user and producer's accuracy values (i.e. minimum omission and commission errors) that could be attained jointly by a low-resolution map. The method has been developed for the specific case of dichotomic classifications and it has been adopted in the evaluation of burned area maps derived from SPOT-VGT with Landsat ETM+ reference data. The use of the Pareto Boundary can help to understand whether the limited accuracy of a low spatial resolution map is given by poor performance of the classification algorithm or by the low resolution of the remotely sensed data, which had been classified.

© 2004 Elsevier Inc. All rights reserved.

Keywords: Omission; Commission; Pareto Boundary

1. Introduction

Thematic land-cover maps are being increasingly produced from image classification of remotely sensed data, which provides consistent and spatially contiguous data over large areas (Foody, 2002). In the last few years, global datasets of low spatial resolution data have been made available to the user community, and consequently several global thematic maps have been produced, such as the International Geosphere–Biosphere Project (IGBP) land-cover map (Loveland et al., 1999), the TRopical Ecosystem Environment observation by Satellite (TREES)

tropical forest maps (Achard et al., 2001, 2002), the Global Land Cover 2000 (GLC2000) global scale vegetation map (Bartholomé et al., 2002) or the Global Burned Area 2000 (GBA2000) burned area map (Grégoire et al., 2003; Tansey et al., in press), the MODIS burned area product (Justice et al., 2002; Roy et al., 2002) or the GLOBSCAR multi-year burned area maps (Simon et al., 2004). The accuracy assessment methods available in the literature relate mainly to mapping investigations on local to regional scales and may not be transferable to coarser scales (Foody, 2002). The development of methods and protocols for the assessment of products derived from moderate to low spatial resolution data is still a research topic, and its need has been recognised as a priority (Chilar, 2000; Foody 2002; Justice et al., 2000; Mayaux & Lambin, 1995).

Mayaux and Lambin (1995) analysed the effects due to the spatial aggregation on the estimation of tropical forest

* Corresponding author. Tel.: +39-0223699457; fax: +39-0223699599.
E-mail addresses: luigi.boschetti@ieec.org (L. Boschetti), stephane@flasseconsulting.net (S.P. Flasse), brivio.pa@irea.cnr.it (P.A. Brivio).

area derived from NOAA-AVHRR. In the analysis of the relationships between area estimates at a fine and a coarse resolution, they indicated the Matheron index and the patch size, both related to the fragmentation of the thematic map, as factors influencing the accuracy of the low-resolution map. More recently, in a systematic analysis of spatial sampling effects on burnt area mapping, Eva and Lambin (1998) underlined the impact of the proportion and spatial pattern of burnt areas at the high spatial resolution as well as of the detection threshold, on the accuracy of burned area estimates from low-resolution satellite data. The study by Smith et al. (2003) on the impact of landscape characteristics on classification accuracy highlighted the impact land-cover heterogeneity and patch size have on classification accuracy, with the probability of a correct classification decreasing with decreasing patch size and increasing heterogeneity.

The comprehensive review of Foody (2002) analyses several aspects of the background and methods currently in use within the remote sensing community and recommended in the research literature. However, no particular mention is given to the problems raised by the accuracy assessment issues of land-cover mapping of large areas from coarse spatial resolution satellite data.

The vast majority of the statistical indices, which are available in the literature for the validation of thematic maps, are derived from the error matrix (Congalton et al., 1983; Story & Congalton, 1986; Stehman (1997) provides a review of these indices. The error matrix is a square matrix which expresses the number of pixels (i.e. the amount of land surface) assigned to a particular land-cover class relative to the actual verified land-cover class. In this article, the columns of the matrix correspond to the reference data, assumed as the truth, and the rows correspond to the classification that is being evaluated. The diagonal elements of the matrix are the correctly classified pixels; therefore, an absolutely correct classification will generate a diagonal matrix.

However, the typical use of the error matrix assumes that reference data and classified data have the same spatial resolution, which is often not the case when evaluating coarse spatial resolution classifications, e.g. data acquired by sensors such as NOAA-AVHRR, TERRA and AQUA-MODIS, SPOT-VEGETATION or AATSR on Envisat. Usually the accuracy assessment or evaluation process is carried out by using samples of either ground data and/or higher spatial resolution multispectral data (such as Landsat-TM or SPOT-HRV). The present work proposes a methodology to analyse the influence of the low spatial resolution of the remotely sensed dataset on the accuracy of the final thematic map products, highlighting the presence of a trade-off between omission and commission errors when employing a hard classification scheme. First the concepts of low-resolution bias and ‘Pareto Boundary’ are introduced. Then this paper presents how the latter concept can be used at different stages of the

classification process, and illustrates its application with some examples.

This work has been developed in the framework of the Global Burned Area 2000 (GBA2000) initiative (Grégoire et al., 2003), coordinated by the Global Vegetation Monitoring Unit of the European Commission Joint Research Centre, aimed at the production of monthly burned area maps at global scale. The maps (<http://www.gvm.jrc.it/fire/gba2000/index.htm>) are derived from low spatial resolution SPOT-VEGETATION data and their accuracy assessment is conducted mainly by comparison with Landsat ETM data, using a multi-level accuracy assessment procedure (Boschetti et al., 2001).

The assessment method described in this paper has been developed specifically for binary classification products (e.g. *dichotomic* classifications where the user is interested in the accuracy of the mapping of one class versus the background). *Dichotomic* classifications are typically used in the study of land-cover disturbances such as flooding or fire effects. Burned area maps will be used in this paper as an illustrative example.

2. Concepts and methods

In the case where a user is interested in the accuracy assessment of a single class, this class often represents a relatively small fraction of a thematic map. In the following, we will refer to generic thematic classes ω_1 and ω_2 , where class ω_1 is the class of interest and ω_2 is always the background. The method could be extended to n classes if the classes are evaluated *one at a time*, merging all the other classes into a generic background. Since the mapping of class ω_1 is the objective of the dichotomic classification itself, the assessment must focus on this class of interest. For this reason, the analysis of the accuracy of low-resolution thematic products will be conducted from the point of view of the ‘non-background’ class ω_1 .

The assessment of thematic product is commonly achieved by comparing it with some samples of reference data assumed to represent the reality. When evaluating global products, a significant number of samples of reference data are required, making the acquisition of ground truth difficult and unrealistic, hence calling on high spatial resolution satellite data and their best possible interpretation. Consequently, the proposed methodology aims at addressing issues linked to the use of samples of high-resolution map as reference data for the accuracy assessment of low-resolution products. Like these product assessments, the presented analysis assumes that reference data represent the reality, i.e. it is free from errors. It has been recognised that, when comparing any spatial datasets, geo-location errors between them may affect heavily the accuracy indices (Foody, 2002; Plourde & Congalton, 2003; Thomlinson et al., 1999). For the case of dichotomic classification, Mayaux and Lambin (1995) proposed a

method based on the spatial aggregation on larger cells (boxes of 11×11 low-resolution pixels) but the compensation for misregistration errors is obtained at the cost of a complete loss of information about the locational accuracy within the boxes. More recently, [Rommel and Perera \(2002\)](#) proposed a method based on the comparison of polygons. However, these methods are not compatible with the use of the error matrix. The Pareto Boundary is based only on the reference data and on the pixel size of the low-resolution data (not on the actual low-resolution map). As a consequence, rather than an issue of misregistration, there is a sub-pixel uncertainty in the positioning of the low-resolution grid on the higher resolution data. The impact of this uncertainty on the Pareto Boundary will be discussed in Section 3.4 and numerically assessed in the two examples of Section 3.5.

Finally, it must be noted that the proposed methodology applies to hard classifiers only; therefore, the techniques developed for *soft* classifiers ([Binaghi et al., 1999](#); [Foody, 1996](#)) will not be taken into account. A soft classification scheme could indeed theoretically achieve 100% accuracy even with mixed pixels, and the computation of the Pareto Boundary actually involves the hardening of an ideal error-free soft classification, as explained in detail in Section 3.1. The discussion of the relative merits and drawbacks of hard and soft classifiers is not the objective of this paper. Instead, the proposed methodology is intended as an

assessment of the cost, in terms of classification errors, of the choice of hard classifiers. To date, the vast majority of small scale (regional, continental or global) land-cover products is constituted by crisp thematic maps. The reasons for such a choice range from the extensive training required by fuzzy classifiers, which could have a cost too high for global maps, to the common practice of using algorithms which are intrinsically hard (like in many of the available burned area mapping techniques) and, last but not least, to the demand of crisp maps by the decision makers and other end-users.

2.1. The low-resolution bias

Let us consider the accuracy assessment of a classified low spatial resolution image, using as reference data a map of the same area produced at high spatial resolution. [Fig. 1](#) illustrates how they could typically relate to each other spatially. While some pixels are entirely classified correctly and a (hopefully limited) number of pixels are totally misclassified, all the remaining pixels have been correctly classified only for a fraction of their extent. Every time that a low-resolution pixel is partially covered by more than one class in the high-resolution map, the accuracy of the low-resolution product will automatically decrease. If the pixel is classified as class ω_1 , then the commission error will increase (because of the fraction of background within the

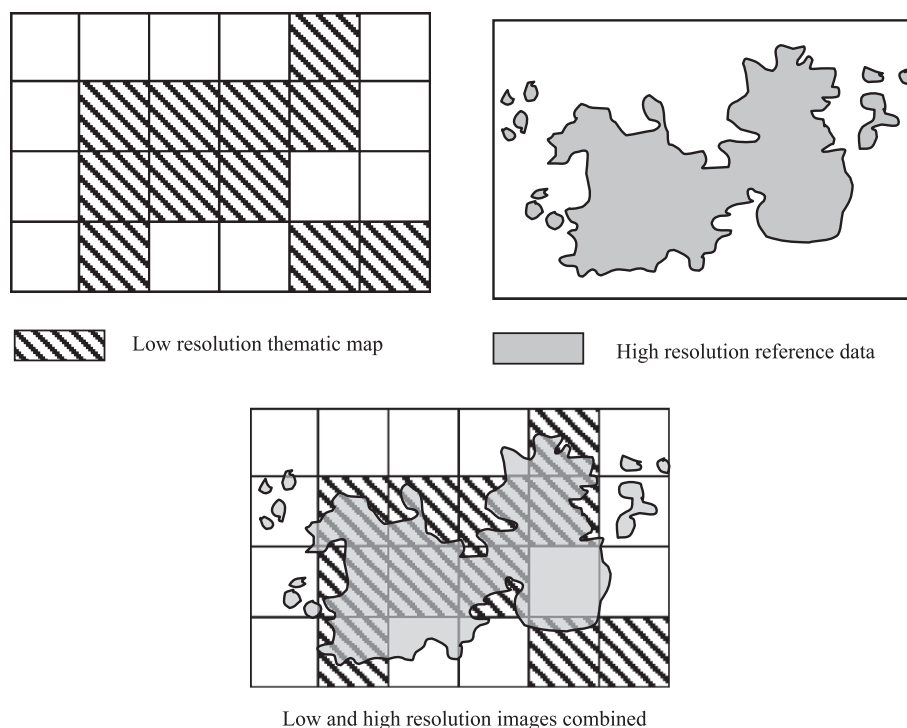


Fig. 1. Two typical dichotomic thematic maps, covering the same area, one derived from low spatial resolution imagery, and the other one from high resolution. At low spatial resolution, the majority of pixels covers a mixed coverage. Consequently, whatever the class they are assigned to, pixels will generate omission or commission errors.

pixel), but if it is classified as background ω_2 , then the omission error increases. A low-resolution pixel covered exactly by 50% class ω_1 and 50% class ω_2 is the extreme case. Whatever label is assigned to it, omission and commission errors will increase by the same amount.

In the situation exemplified in Fig. 1, it is possible to identify intuitively, pixel by pixel, whether the error (either commission or omission) is due to an actual misclassification, or whether it is due to mixed coverage of the pixel. The error matrix does not take into account any of such ‘contextual intuitive’ consideration. It has indeed already been recognised that the accuracy measures derived from the error matrix, which assume implicitly that the testing samples are pure, are often inappropriate for the assessment of remotely sensed datasets (Binaghi et al., 1999; Foody, 1996, 2002).

It is important to note that the number of low-resolution pixels with mixed land cover is linked to the intrinsic characteristic of the features on the ground (i.e. in the reference data) and it is a function of their shape, size and fragmentation (Eva & Lambin, 1998; Mayaux & Lambin, 1995; Woodcock & Strahler, 1987). Henceforth we refer to the inaccuracy introduced by the difference in spatial resolution between high and low resolution data as the ‘low-resolution bias’.

A quantitative analysis of the impact of this effect on the accuracy of the low spatial resolution thematic maps could benefit several phases of the classification process. Given that the low-resolution bias depends only on the landscape heterogeneity of the class of interest in the high-resolution reference data and on the pixel size of the low-resolution data, an a priori feasibility analysis could be carried out, using a sample of high-resolution reference data, for understanding the implications of the use of a low spatial resolution product, even before embarking in the classification processing of the low-resolution dataset. Such a quantitative analysis could also be useful to the analyst during the assessment of the first results of a classification algorithm, or when choosing among different available algorithms because it allows understanding how much of the error is due to the algorithm (and could be avoided by using a better one) as well as how much is due to the spatial characteristics of the low-resolution data itself.

2.2. Omission and commission errors as conflicting objectives

The main effect of the low-resolution bias concept introduced above is that the reduction of the omission and the reduction of the commission errors become two partially conflicting objectives. Once all pixels entirely belonging to either class have been classified, the residual errors shown by the error matrix cannot be avoided in any way. Any reduction in commission error will be attained at the cost of increasing the omission error and vice versa,

because it will imply a change of label of at least one mixed pixel.

Table 1 is an example of confusion matrix. Diagonal elements represent the pixels correctly classified, while off-diagonal elements represent errors, either of commission or omission (Congalton, 1991). From the error matrix, it is possible to derive two class specific indices, Commission Error (Ce) and Omission Error (Oe), which quantify the classification errors of commission and omission, respectively. For class ω_1 , these are formulated as follows:

$$Ce = \frac{p_{12}}{p_{11} + p_{12}} \quad (1)$$

$$Oe = \frac{p_{21}}{p_{11} + p_{21}} \quad (2)$$

The Commission Error Ce of a generic class ω_1 is the percentage of pixels classified as class ω_1 which do not belong to that class according to the reference data (commission), while the Omission Error Oe is the percentage of the pixels, belonging to class ω_1 in the reference data, which have not been classified as such (omission). Omission and commission errors make excellent candidate *indices* to represent the situation of reducing omission and commission errors as conflicting objectives. In the following, these two indices are used in order to model the low-resolution bias problem as a minimisation problem, i.e. given the sensor’s spatial resolution, what is the minimum level of omission and commission errors that any thematic map of given spatial resolution will have.

It should be noted that Omission Error and Commission Error are linked to the widely used User’s Accuracy and Producer’s Accuracy (Congalton & Green, 1999) by the relations:

$$Ce = \frac{p_{12}}{p_{11} + p_{12}} = \frac{p_{12} + p_{11} - p_{11}}{p_{11} + p_{12}} = 1 - \frac{p_{11}}{p_{11} + p_{12}} = 1 - Ua \quad (3)$$

$$Oe = \frac{p_{21}}{p_{11} + p_{21}} = \frac{p_{21} + p_{11} - p_{11}}{p_{11} + p_{21}} = 1 - \frac{p_{11}}{p_{11} + p_{21}} = 1 - Pa \quad (4)$$

2.3. The Pareto optimum

When in the process of developing algorithms for the creation of thematic products, such as regional, continental

Table 1
Typical confusion matrix for a dichotomic thematic product

Classified data	Reference data	
	Class ω_1	Class ω_2 (background)
Class ω_1	p_{11}	p_{12}
Class ω_2 (background)	p_{21}	p_{22}

or global burned area maps, it is important to have some tools to quantify as representatively as possible their performance, independently to the limitation of the actual low spatial resolution data.

As discussed above, the accuracy indices alone are not sufficient to evaluate how well a classifier has performed, because at a certain point, it is not possible to minimise both omission and commission errors at the same time. To rank the classification results and to establish when a thematic map is better than another, we can borrow the *Pareto Boundary of efficient solutions* from Economics Theory, where the so-called ‘Pareto optimum’ (Pareto, 1906) is largely used in the case of decisions with multiple objectives (Keeney & Raiffa, 1976). According to Pareto’s definition of efficiency, the optimum allocation of the resources is attained when it is not possible to make at least one individual better off while keeping others as well off as before. ‘Efficiency’ does not imply any evaluation of the ‘equity’ or ‘fairness’ of the distribution, and that it is possible to find up to an infinity of efficient solutions.

In our case of image classification, as ranking criterion, we assume that preference will be given to lower commission or omission error; i.e. that *ceteris paribus* a classified image with lower omission (commission) will be preferred to one with higher omission (commission). Hence, we can say that classified image A *dominates* (Keeney & Raiffa, 1976) classified image B whenever

$$\begin{cases} Oe(A) < Oe(B) \\ Ce(A) \leq Ce(B) \end{cases} \quad (5i)$$

or

$$\begin{cases} Oe(A) \leq Oe(B) \\ Ce(A) < Ce(B) \end{cases} \quad (5ii)$$

With these assumptions, we can easily adapt Pareto’s definition to our case, by saying that a low-resolution thematic product is *Pareto optimal* if it is impossible to find another classification with less commission error (omission error) with at the most the same omission error (commission error). The set of Pareto optimal classified images will dominate any other possible classification and it is known as the *Pareto Boundary* (also known as the *Efficient Frontier* or the *Pareto optimal set*).

Adopting the Pareto efficiency approach, it follows that:

- we are interested only in the ordinal character of the accuracy indices, not in the cardinal character, e.g. given three thematic maps A, B and C, with omission error

$Oe(A)=0.2$, $Oe(B)=0.3$, $Oe(C)=0.6$, we are interested in the relationship that $Oe(A) < Oe(B) < Oe(C)$, regardless of the actual values of $Oe(A)$, $Oe(B)$ and $Oe(C)$. Any monotonic function of omission and commission errors would lead to the same optimal set of classification maps.

- no preference structure other than the relationship of dominance is expressed, i.e. no preference can be expressed among a set of classification maps if none of them is dominating any other. Further ranking (Section 3.3) could therefore only be based on the fairness of a classification (i.e., to what extent it is better to have omission rather than commission error or omission rather than commission error) but there can be no fixed rule to evaluate it. The user needs to specify if there is a greater cost associated to omission errors than to commission errors.

As an example of the difference between efficiency and fairness, let us consider a low-resolution map where there are no class ω_1 pure pixels, but many pixels with 99% class ω_1 . The classification with complete omission (and, of course, no commission error) of ω_1 will be an *efficient* solution, because to reduce the omission error, a commission error, even if minimal, must be introduced. At the same time, for most of the users such a classification will be almost useless, because class ω_1 has been completely missed. In this sense, we will say that for these users, the classification is not *fair*. Anyway, while the efficiency of the classification can be evaluated based only on intrinsic characteristics of the area mapped (the landscape heterogeneity), the fairness can be evaluated only from a specific user’s point of view, and when the user’s preferences are known.

2.4. Defining the set of Pareto-optimal classifications

A set of Pareto-optimal classified images, with the properties defined in the previous paragraph, could be obtained starting from the high-resolution reference map (our ‘true’ data). Each *optimal* classification will have a pair of error rates (Oe, Ce), that uniquely identify a point in the omission/commission $Oe/Ce = \{(Oe, Ce) \in \mathbb{R} \times \mathbb{R} \mid Oe \in [0, 1], Ce \in [0, 1]\}$ bi-dimensional space (henceforth noted Oe/Ce). The line joining all these points is the Pareto Boundary related to the specific high-resolution reference map and to a specific low-resolution pixel size. As all the points belonging to the Pareto Boundary represent the error level of efficient classification, by the definition of efficiency given in the previous paragraph, it follows that the Pareto Boundary divides the Oe/Ce in two regions. Any possible low-resolution thematic products can be associated only to the points in the region above or, at the most, on the Pareto Boundary. The point (0,0), with neither omission nor commission error, will unfortunately lie in the unreachable region!

The terms ‘efficient’ or ‘optimal’ will be used as synonyms and will always mean efficient (optimal) under Pareto’s definition. The Pareto-dominance will be adopted

as the main ranking criterion, i.e., ‘A performs better than B’ is equivalent to say ‘A dominates B’.

3. Using the Pareto Boundary

3.1. Calculating the Pareto Boundary

To create the boundary, we only need to use:

- (i) the high-resolution reference map,
- (ii) the *size* of the low-resolution cell (not the actual low-resolution thematic product).

Fig. 2 exemplifies how it is possible to determine point by point the Pareto Boundary and to eventually trace it in the ‘Oe/Ce’ bi-dimensional space. We start by applying the low-resolution grid to the high-resolution reference map, assigning to each new low-resolution cell a value corresponding to the percentage of class ω_1 it covers (Fig. 2a and b). This new low-resolution cell “product” could be seen as a sort of ideal *soft* low-resolution classification. Defining a threshold t on the percentage of class ω_1 present within a low-resolution cell, above which it would be selected as ω_1 , we are actually creating “ideal-virtual” *hard* classification maps for each threshold value. We have thus an optimal classification, because it is possible to reduce the commission error only by raising the threshold t and consequently increasing the omission error, and vice-versa. Consequently, a discrete set of t values covering the range (0, 1] will generate a set of Pareto-optimal *hard* classifications (Fig. 2c): omission and commission error pairs computed for each threshold value will identify a set of points in Oe/Ce. The line linking these points will be a discrete approximation of the Pareto Boundary (Fig. 2d).

More rigorously, once the soft low-resolution map with low-resolution cell size L has been obtained, commission

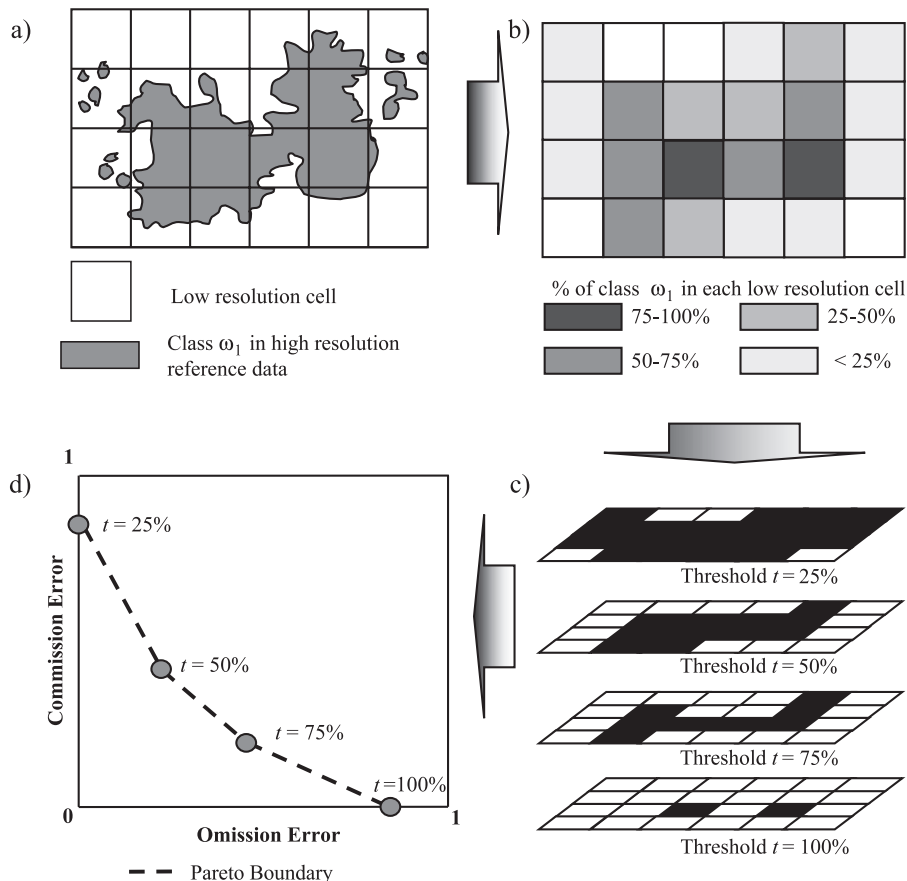


Fig. 2. The procedure for generating a discrete set of points belonging to the Pareto Boundary, starting from the high-resolution map, for a desired low spatial resolution. (a) A low-resolution grid is overlaid on the high-resolution reference map. (b) The percentage of class ω_1 is computed for each cell. (c) A set of low-resolution products is generated by thresholding the percentage of class ω_1 ; threshold t varies in the interval (0, 1]. All of these maps are *efficient solutions* according to Pareto's criterion. The map generated with $t = 100\%$ will have no commission (no mixed pixels included) but large omission, and the map with $t = 1\%$ will have no omission (all the mixed pixels are included) but large commission; in general, the higher t , the lower the commission and the higher the omission. (d) The confusion matrix is produced for each one of these maps. Omission error and commission error are derived and plotted in the omission error/commission error space. As all the maps are *efficient solutions*, the line linking the corresponding points is a discretisation of the Pareto Boundary.

and omission error can be easily computed as a function of the threshold value t :

$$A_L(t) = \int_t^1 N_{Li} di \quad (6)$$

$$A_H = \int_{0+}^1 i N_{Li} di \quad (7)$$

$$O_L(t) = \int_{0+}^t i N_{Li} di \quad (8)$$

$$C_L(t) = \int_t^1 (1 - i) N_{Li} di \quad (9)$$

$$Ce(t) = C_L(t)/A_L(t) \quad (10)$$

$$Oe(t) = O_L(t)/A_H. \quad (11)$$

In the discrete case, Eqs. (6)–(9) can be written as:

$$A_L(t) = \sum_{i=t}^1 N_{Li} \quad (12)$$

$$A_H = \sum_{i=t}^1 i N_{Li} \quad (13)$$

$$O_L(t) = \sum_{i>0}^t i N_{Li} \quad (14)$$

$$C_L(t) = \sum_{i=t}^1 (1 - i) N_{Li}. \quad (15)$$

Where N_{Li} is the number of cells with fraction i ($0 < i < 1$) coverage of class ω_1 . $A_L(t)$ is the area of class ω_1 estimated at low spatial resolution, with the threshold t . A_H is the area of class ω_1 in the high spatial resolution reference data. $O_L(t)$ is the area omitted by the optimal low-resolution classification obtained with threshold t . $C_L(t)$ is the area committed by the optimal low-resolution classification obtained with threshold t . $Ce(t)$ is the commission error. $Oe(t)$ is the omission error.

All the areas are expressed in low-resolution cell size units.

From the definition of Pareto-efficiency, it follows that a classification is Pareto-optimal only if all the pure pixels have been correctly classified. Let us consider a classification, where a pure pixel belonging to class ω_1 has been assigned to class ω_2 . By changing the label of that pixel only, the commission error would be reduced with no cost in terms of omission error. This contradicts Pareto's definition

of efficiency and it follows that at the beginning, the classification was not optimal.

Although the proposed method for the determination of the Pareto Boundary applies only to dichotomic classifications, it is possible to force it to the general case of n different classes. The analysis has to be repeated for each class, keeping the class of interest as ω_1 and collapsing all the other classes into the background ω_2 . The result will be a set of n plots on the Oe/Ce space, each showing the Pareto Boundary of one single class.

3.2. Interpreting the Pareto Boundary

Given a generic low-resolution thematic product, obtained via a classification procedure, the confusion matrix can be calculated and it is possible to plot the position of the corresponding accuracy measures of omission and commission error on Oe/Ce space. Let us call A the point of the bi-dimensional space, which is associated with this low-resolution map. With this accuracy point A in the space, using the relationship of *dominance* introduced above, it is possible to partition the plane into four regions, as exemplified in Fig. 3. The region at the upper right of the point A is the region where commission and omission errors are higher




Fig. 3. The partitioning of the Oe/Ce space. The Pareto Boundary divides it into two regions, reachable/unreachable. Given a low-resolution product accuracy A with accuracy indices (Oe_A , Ce_A), the corresponding point will introduce four more regions: a region dominating A (where both omission and commission are lower than Oe_A and Ce_A), a region dominated by A (where both omission and commission are higher than Oe_A and Ce_A) and two regions of indifference (where either omission or commission are lower than Oe_A and Ce_A).

than (Oe_A, Ce_A) , and consequently, it is *dominated* by A. The region at the lower left is the region where omission and commission are lower than (Oe_A, Ce_A) and consequently it *dominates* A; the lower limit of this region is the Pareto Boundary itself. The other two (lower right and upper left) are the regions of *indifference*, where either omission or commission error are lower than Oe_A and Ce_A .

In the case of Fig. 3, the accuracy of the thematic product corresponding to point A is not very high (both Oe and Ce are about 0.4). Using its position with respect to the Pareto Boundary, it is possible to determine how close the actual map comes to meeting the optimal conditions represented by the boundary itself. At the most, if the point is on the boundary, all the errors will be a result of the difference of spatial resolution (the low-resolution bias) more than of the algorithm's weakness.

Fig. 4 shows the situation when more classification methods are applied to the same low-resolution image: points A, B and C all correspond to non-optimal thematic products, while points P_1 , P_2 , P_3 and P_4 belong to the Pareto Boundary and therefore correspond to optimal thematic products. Here the relationship of dominance can be used as a criterion for ranking the different classification procedures: both A and B dominate C, performing better both for omission and for commission. By definition, the points on the boundary cannot be dominated by any other point, but this fact does not mean that an optimal solution is always to be preferred to a non-optimal one. For instance, if the choice is between A and a classifier whose performance is P_4 , among which there is no relationship of dominance, the selection should depend on the specific user's need. Although not being an efficient solution, which means that further improvement is possible, A has lower omission error than the efficient solution P_4 . As P_4 has no commission error

error, but a drastic omission error, only those users, for whom only omission error is the critical issue, would choose a classification procedure leading to P_4 , as the more suitable option.

3.3. The user's cost function

So far, the analysis and the ranking of points on the Oe/Ce bi-dimensional space have been conducted solely from the point of view of the Pareto optimum. However, as highlighted in Section 2.3, a Pareto-optimal classification is not necessarily a desirable one from the user's point of view: intuitively, not all the points of the Pareto Boundary are alike. For instance, the intersection between the boundary and the x axis in Fig. 4 represents an efficient classification, with no commission error but a dramatic level of omission error and few users would actually accept such a map.

From the assumptions made in Section 2.2, it is also clear that, in the regions of Pareto-indifference, ranking can be based on the fairness of the classification, and that the fairness depends on the specific user's needs.

The idea of introducing the monetary cost of classification errors in the error matrix has been considered by Smits et al. (1999) for the case of land use maps for allocating agricultural subsidies. The notion of *cost* could be extended to include all the non-monetary consequences of misclassification.

Assuming that the cost function for the user is a function of Ce and Oe , it will be possible to associate a value $C = C(Ce, Oe)$ to all the points of the Ce/Oe space and even rank in order of preference the points of the Pareto Boundary. In order to do that, we will adapt to the classification problem the 'indifference curves' approach for structuring preferences and value functions reported by Keeney and Raiffa (1976).

The underlying assumptions that need to be made are:

- (i) the cost (either monetary or non-monetary) for the user is a function of Ce and Oe only,
- (ii) the cost is monotonically increasing both in Ce and Oe .

Typical shapes of cost function are linear or hyperbolic combinations of Ce and Oe :

$$C(x) = c_1 Ce + c_2 Oe \quad (16i)$$

$$C(x) = 1 - [(1 - Ce)^\alpha (1 - Oe)^\beta] \quad (16ii)$$

which can be combined in a general form:

$$C(x) = c_1 Ce + c_2 Oe + c_3 \{1 - [(1 - Ce)^\alpha (1 - Oe)^\beta]\} \quad (16iii)$$

where c_1 , c_2 , c_3 , α and β are the parameters which determine the sensitivity of the function to commission or omission and need to be set on the basis of the requirements of the user.

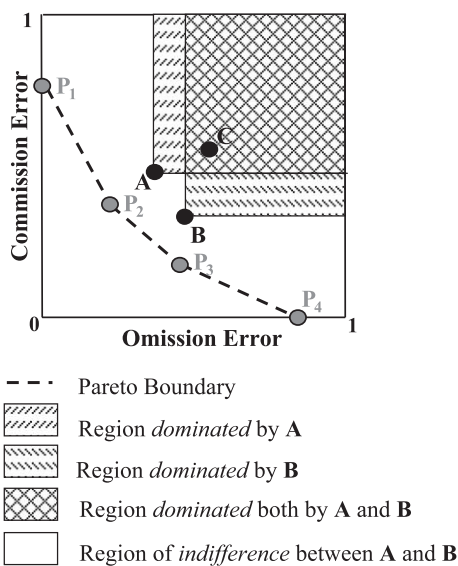


Fig. 4. The relationship of Pareto-dominance as a criterion for ranking the performance of different classification algorithms of the same low-resolution data. Points A and B in the figure dominate point C.

The higher are c_1 and α , the more the user is worried by commission error, the higher are c_2 and β , the more the user would like to avoid omission error. Fig. 5 shows an example of linear and hyperbolic iso-lines in the C_e/O_e space.

When the user bears a monetary cost for the classification errors, the cost function can be easily determined; in all the other cases, there is no way to determine unequivocally the coefficients or even the general form of the cost function. Consequently, the whole idea of cost function must be taken as a way of formalising in a rigorous way an analysis that is substantially a qualitative one.

Associating the concept of Pareto-Dominance and Pareto Boundary (Section 2.3) with the cost function defined above, it is possible to solve the problems of incomplete ranking, highlighted in Section 3.2, by observing that:

- (i) A dominates $B \Rightarrow C(A) < C(B)$, i.e. whenever A dominates B, the cost of A is always lower than the cost of B.
- (ii) The criterion of minimum cost can be used for ranking in the regions of Pareto-indifference.
- (iii) In case of Pareto-indifference and equal cost, the point with the smaller Euclidean distance from the Pareto Boundary is preferred.

As the Pareto Boundary is a region of indifference (all the points are efficient, thus by definition not dominated), the points of the Pareto Boundary can be ranked according to the associated cost: the point of tangency between the Boundary and the iso-lines is the ideal point representing the best achievable result, for the specific user.

Fig. 6 shows an example of complete ranking of the results of different methods in the C_e/O_e space:

- Point P_2 is the point representing the best possible method for the specific user.

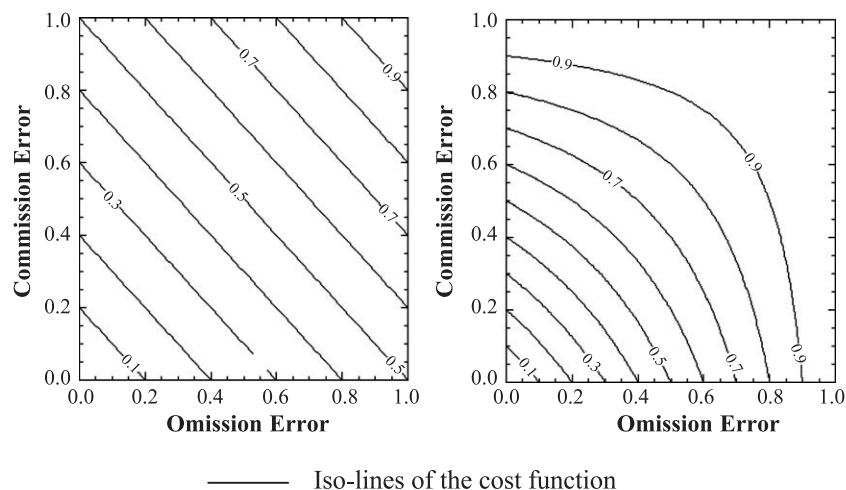


Fig. 5. Example of iso-lines of linear and hyperbolic cost functions.

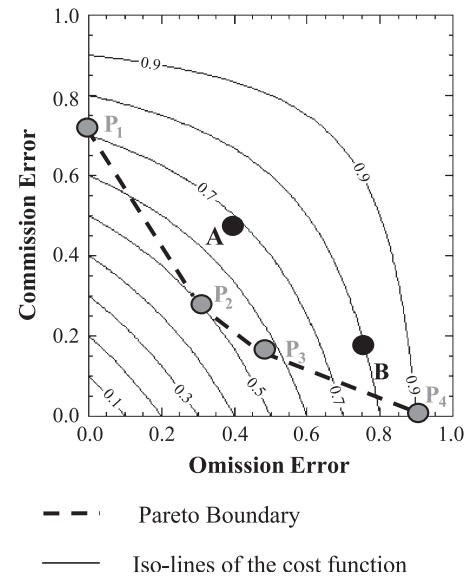


Fig. 6. Ranking combining the Pareto-dominance and the user's cost function. Point P_2 of tangency between the Pareto Boundary and the iso-cost lines is the point with lowest misclassification cost, compatible with the low-resolution bias. Points A and B are ranked according to the cost function.

- A is neither dominating nor dominated by B, therefore, the ranking is based on the cost function. Because $C(A) < C(B)$, A should be selected as more respondent to the requirements of the user.
- The difference $C(A) - C(P_2)$ is an index of the margin of improvement of the classifier.

Lark (1995) introduced the concept that the classification algorithm can be modified to emphasize whichever is more important, user's or producer's accuracy for a particular user. Point P_2 could be seen as the target of such an adjustment of the classification method.

3.4. The misregistration issue: A stochastic approach

The Pareto Boundary, being computed from the high-resolution data alone, is not affected by misregistration between the low-resolution map and the reference data. As it requires that the a low-resolution grid is superimposed on the high-resolution data, there is an uncertainty of ± 0.5 low-resolution pixel units both in x and in y directions and this uncertainty will propagate on the positioning of the Pareto Boundary points in the Oe/Ce space. Being the Pareto Boundary, a function of the pattern of class ω_1 in the specific images, such an uncertainty cannot be determined analytically, but needs to be assessed numerically, generating a random sample of shifts (Δx , Δy) within the range $[-0.5, 0.5]$ and computing the variance of the omission and commission error pair values of all the points belonging to the Boundary. As a result, the Pareto Boundary is no longer a deterministic set of points (a line) but becomes a region of the Oe/Ce space, described by a probability distribution.

When the Pareto Boundary is used, as described in Sections 3.2 and 3.3, to evaluate low-resolution classification results, the misregistration between low and high-resolution data is an issue that should be taken into account, because it will affect the omission and commission error of the classification (Foody, 2002; Plourde & Congalton, 2003; Thomlinson et al., 1999). Automated co-registration software tools based on correlation algorithms (Baltasvias, 1991; Carmona-Moreno, in press) provide sub-pixels accuracy, set aside the case of strongly nonlinear deformations. The coregistration error could be seen as an uncertainty and, as it is possible to estimate its variance, it is also possible to estimate its propagation on the position of the point in Oe/Ce , using the same numerical procedure adopted for the points of the Pareto Boundary.

Under the usual assumptions of the error matrix (same spatial resolution of the datasets, no misregistration errors), the problem of minimising the classification error is a mono-dimensional, deterministic one. By taking into account the low-resolution bias, the problem is still deterministic, but requires the use of a bi-dimensional space; by taking also into account the misregistration issue, the problem becomes a bi-dimensional, stochastic one.

3.5. Examples: Application of the Pareto Boundary to burned area maps

As discussed in Section 3.1, the Pareto Boundary can be calculated for any low-resolution pixel size starting from the high-resolution reference data alone. Images with different levels of fragmentation will generate different Pareto boundaries. In general, the higher the fragmentation, the lower the accuracy that can be achieved and the higher the proportion of the Oe/Ce space that is not reachable. This is illustrated in Fig. 7, where the Pareto Boundary is computed for two burned area maps derived from Landsat7 ETM data, by

means of a maximum likelihood supervised classification (Boschetti et al., 2002; Brivio et al., 2003). The two images, respectively, scene 179-74 (Fig. 7a and b) acquired over Namibia on October 31st, 2000, and scene 177-57 (Fig. 7c and d), acquired on December 2nd, 1999 over Central African Republic and Democratic Republic of Congo, were part of the reference dataset used for the accuracy assessment of continental burned area maps obtained from 1-km spatial resolution SPOT-VGT data. Continental burned area maps are a situation where relatively frequent mapping is required over a very large area. The ETM+ scenes used for the validation of the low-resolution data are a sample, both in the space and in the time domain, as it would be impossible (not only because of the cost, but also because of insufficient data availability) to systematically use high resolution data for the production of multitemporal global thematic maps. As explained above, the Pareto Boundary is calculated for a given low spatial resolution grid, 1 km \times 1 km in this example. Fig. 7a and c present the fraction of observed burned area, in each 1 km \times 1 km pixel (what was referred to as *soft classification* in Section 3.1 above); and Fig. 7b and d) the discrete Pareto Boundary. The higher the number of 1-km pixels with a low percentage of burned area, the stronger the effect of the *low-resolution bias* and the lower the accuracy that can be attained by a low-resolution classification. The two scenes show a very different level of heterogeneity. In scene 179-74 (Namibia), 71% of the high-resolution burned pixels fall into low-resolution cells with a fraction of 50% burned area or more, while in scene 177-57 (Central Africa), the cells with a fraction of 50% burned area or more account only for 27% of the total high-resolution burned pixels. To illustrate the fragmentation, we use the Matheron index, identified by Eva and Lambin (1998) as one of the indicators of landscape fragmentation, which in the case of two classes can be computed as:

$$M(\omega_1) = \frac{p\omega_1}{\sqrt{A\omega_1}\sqrt{Atot}} \quad (17)$$

where $p\omega_1$ = total perimeter of class ω_1 , $A\omega_1$ = area of class ω_1 , $Atot$ = total area.

The Matheron Index for the Namibia scene is 4.6, while it is 6.2 for the more fragmented Central Africa scene. The different level of fragmentation of the two thematic maps is reflected in the respective Pareto Boundary plots. The Pareto Boundary of the more fragmented scene of Central Africa is always above that of the Namibia scene and its unreachable region in the Oe/Ce space is considerably larger. As a consequence, the accuracy levels that can be attained by any burned area map of the region covered by scene 177-57 are considerably smaller than that of the other scene. For instance, a classification with $Ce=0.42$ and $Oe=0.43$, which could be considered quite a poor performance, is in fact on the Pareto Boundary, meaning that such level of error is entirely due to the low-resolution bias.

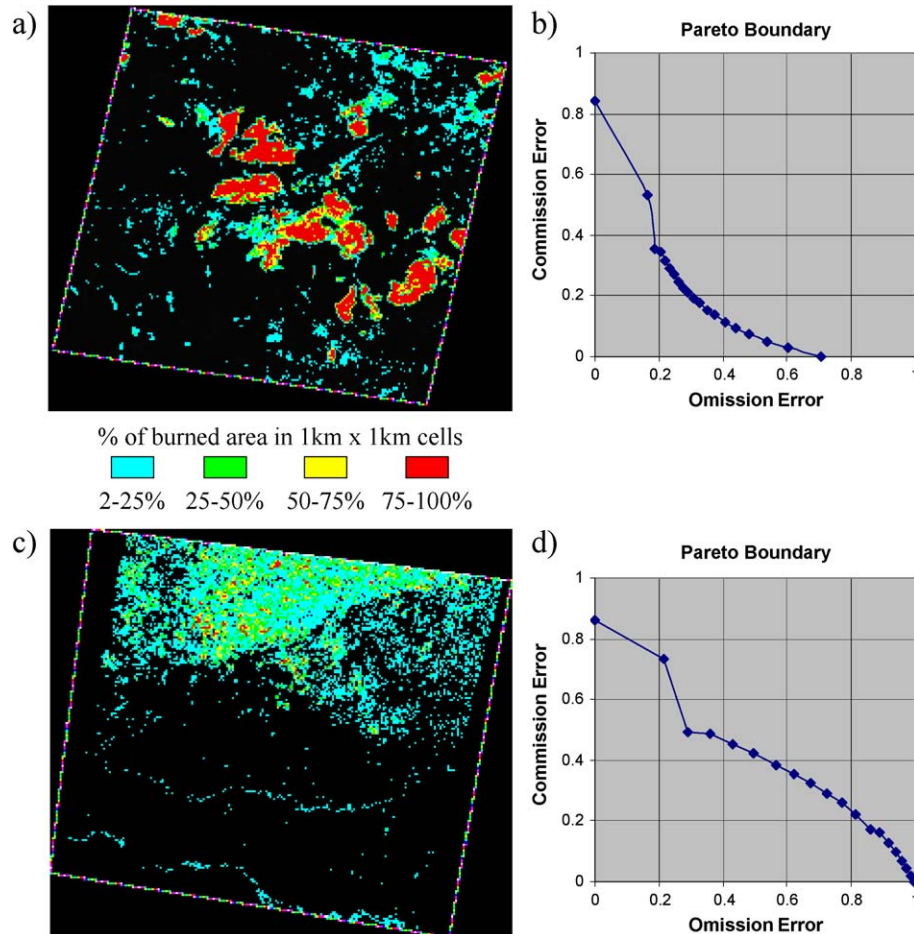


Fig. 7. Two examples of Pareto Boundary calculated for a grid of 1×1 -km resolution with reference burned area maps derived from Landsat7-ETM data, scene 179-74 of October 31st, 2000 (a and b) and scene 177-57 of December 2nd, 1999 (c and d). The reference maps have been produced by means of a maximum likelihood supervised classification. On the left, the two reference maps degraded to a spatial resolution of $1 \text{ km} \times 1 \text{ km}$ with displayed the percentage of burned area within each pixel. On the right, the corresponding Pareto Boundary.

The variance of the Pareto Boundary, due to the uncertainty in positioning the low-resolution grid, has been computed for both datasets. A random sample of 100 shift values ($\Delta x, \Delta y$) has been generated, under the hypothesis of uniform distribution. The Pareto Boundary has been computed for each couple of values. As the standard deviation of O_e and C_e is always below 0.01, it is possible to conclude that the uncertainty in the positioning of the Boundary is negligible.

4. Conclusions

The accuracy assessment techniques, based on the error matrix, have been traditionally applied to reference data and classified data with the same spatial resolution. When thematic maps derived from large datasets of low-resolution remotely sensed data (e.g. NOAA-AVHRR, SPOT-VGT or MODIS) are validated with higher resolution reference data, such as Landsat TM or SPOT-HRV, both simple theoretical analysis and empirical results indicate that the class specific accuracy indices derived from the error matrix

are not adequate when low-resolution pixels in the reference data correspond to more than one class. Similarly, low-resolution thematic maps, often created at regional, continental and global scale, cannot achieve 100% accuracy (*low-resolution bias*).

For the specific case of dichotomic classifications, the low-resolution bias generates a conflict between omission and commission errors. This paper introduced the concept of 'Pareto-dominance', widely used in Economics and in Decision Analysis, to analyse quantitatively the influence of the low-resolution bias and to determine the set of maximum user's and producer's accuracy values that can be attained by any possible thematic product at given low spatial resolution.

A numerical procedure to compute the Pareto Boundary on the Commission Error/Omission Error bi-dimensional space is provided, requiring only the higher resolution reference data and the spatial resolution of the low-resolution thematic product.

The Pareto Boundary allows the definition of an upper limit to the accuracy of the possible solutions, which

delimits a region of unreachable accuracy. All the points of the line are equally efficient from the point of view of omission and commission errors, as they represent possible classifications where all the errors are due solely to the low-resolution bias. Nonetheless, if the preference function of the user of the specific map is known, it is possible to derive from the accuracy figures a cost function that can be combined with the Pareto Boundary to identify the point of the boundary which has the lowest cost for the user. This point of the Commission Error/Omission Error space represents the accuracy figures of the best map that could be theoretically achieved, according to the specific user's needs. The difference between the cost of such a point and the cost derived from the accuracy figures of a given thematic map is an index of the performance of the classification algorithm adopted.

Being computed only from high-resolution data, the Pareto Boundary can be easily used as a tool for understanding the implications of the spatial resolution in preliminary feasibility studies. Using a sample of high-resolution data over a region of interest, documentation of the low-resolution bias could help identify the relevance (required accuracy) of a low-resolution product for the target and region of interest, even before starting the development of algorithms. In case of regional or continental scale low-resolution thematic products, the computation of the Pareto Boundary could be calculated for a selection of sites, in order to characterise the low-resolution bias for the relevant ecosystems.

The proposed methodology based on the Pareto Boundary could be also used by the analyst during the classification process. Typically, after the first phase of implementation of a classification method and when the first results are available, the analyst needs to decide whether the classification is accurate enough, or the classifier needs to be improved. If a sample of reference data is available, quantifying the low-resolution bias allows understanding whether the errors described by the error matrix could be reduced using a better algorithm, or are due only to the low-resolution bias and therefore depend on the intrinsic characteristics of the area mapped. Furthermore, when more than one method is available and a choice has to be made, combining the Pareto Boundary with the analysis of the preferences of the user, it is possible to rank the methods and to select the one which corresponds better to the specific user's needs.

Acknowledgements

The authors would like to acknowledge the contribution of Marta Maggi, who kindly provided one of the TM classifications used as examples, and to thank Hugh Eva and Jean Marie Grégoire of JRC, for the useful comments on the original manuscript; we are also grateful for the comments from the anonymous reviewers, which helped to strengthen and clarify the paper.

The first author is supported by a research fellowship from the European Commission-Directorate General JRC.

References

- Achard, F., Eva, H., & Mayaux, P. (2001). Tropical forest mapping from coarse spatial resolution satellite data: Production and accuracy assessment issues. *International Journal of Remote Sensing*, 22(14), 2741–2762.
- Achard, F., Eva, H., Stibig, H. J., Mayaux, P., Gallego, J., Richards, T., & Malingreau, J. P. (2002). Determination of deforestation rates of the world's humid tropical forests. *Science*, 297, 999–1002.
- Baltsavias, E. P. (1991). Geometrically constrained multiphoto matching. *Mitteilungen*, vol. 49. Zurich: Institute of Geodesy and Photogrammetry, ETH.
- Bartholomé, E., Belward, A. S., Achard, F., Bartalev, S., Carmona-Moreno, C., Eva, H., Fritz, S., Grégoire, J.-M., Mayaux, P., & Stibig, H.-J. (2002). *GLC 2000-Global Land Cover mapping for the year 2000-Project status November 2002*. Ispra, Italy: Publication of the European Commission, JRC, EUR 20524 EN.
- Binaghi, E., Brivio, P. A., Ghezzi, P., & Rampini, A. (1999). A fuzzy set-based accuracy assessment of soft classifications. *Pattern Recognition Letters*, 20, 935–948.
- Boschetti, L., Flasse, S., Jacques de Dixmude, A., & Trigg, S. (2002). A multitemporal change-detection algorithm for the monitoring of burnt areas with SPOT-Vegetation data. In L. Bruzzone, & P. Smith (Eds.), *Analysis of multi-temporal remote sensing images* (pp. 75–82). Singapore: World Scientific Publishing.
- Boschetti, L., Flasse, S., Trigg, S., Brivio, P. A., & Maggi, M. (2001). A methodology for the validation of low resolution remotely sensed data products. *Proceedings of 5th ASITA conference (Rimini 9–12 October 2001) vol. 1* (pp. 293–298). Asita, Milano.
- Brivio, P. A., Maggi, M., Binaghi, E., & Gallo, I. (2003). Mapping burned surfaces in Sub-Saharan Africa based on multi-temporal neural classification. *International Journal of Remote Sensing*, 20, 4003–4016.
- Carmona-Moreno, C. (2004). Vegetation L-band processing system. Geometric performance and spatio-temporal stability. *International Journal of Remote Sensing* (in press).
- Chilar, J. (2000). Land cover mapping of large areas from satellites: Status and research priorities. *International Journal of Remote Sensing*, 21, 1093–1114.
- Congalton, R. G. (1991). A review of assessing the accuracy of classifications of remotely sensed data. *Remote Sensing of Environment*, 37, 35–46.
- Congalton, R. G., & Green, K. (1999). *Assessing the accuracy of remotely sensed data* (p. 137). Florida: CRC Press.
- Congalton, R. G., Oderwald, R. G., & Mead, R. A. (1983). Assessing Landsat classification accuracy using discrete multivariate analysis statistical techniques. *Photogrammetric Engineering and Remote Sensing*, 49, 1671–1678.
- Eva, H., & Lambin, E. F. (1998). Remote sensing of biomass burning in tropical regions: Sampling issues and multisensor approach. *Remote Sensing of Environment*, 64, 292–315.
- Foody, G. M. (1996). Approaches for the production and evaluation of fuzzy land cover classification from remotely sensed data. *International Journal of Remote Sensing*, 17, 1317–1340.
- Foody, G. M. (2002). Status of land cover accuracy assessment. *Remote Sensing of Environment*, 80, 185–201.
- Grégoire, J.-M., Tansey, K., & Silva, J. M. N. (2003). The GBA2000 initiative: Developing a global burned area database from SPOT-VEGETATION imagery. *International Journal of Remote Sensing*, 24, 1369–1376.
- Justice, C., Belward, A., Morissette, J., Lewis, P., Privette, J., & Baret, F. (2000). Developments in the validation of satellite sensor products for

- the study of land surface. *International Journal of Remote Sensing*, 21(17), 3383–3390.
- Justice, C. O., Giglio, L., Korontzi, S., Owens, J., Morisette, J. T., Roy, D., Descloitres, J., Alleaume, S., Petitcolin, F., & Kaufman, Y. (2002). The MODIS fire products. *Remote Sensing of Environment*, 83, 244–262.
- Keeney, R. L., & Raiffa, H. (1976). *Decisions with multiple objectives: Preferences and value tradeoffs*. New York: Wiley.
- Lark, R. M. (1995). Components of accuracy of maps with special reference to discriminant analysis on remote sensor data. *International Journal of Remote Sensing*, 16(8), 1461–1480.
- Loveland, T. R., Zhu, Z., Ohlen, D. O., Brown, J. F., Reed, B. C., & Yang, L. (1999). An analysis of the IGBP Global Land-Cover characterization process. *Photogrammetric Engineering and Remote Sensing*, 65, 1021–1032.
- Mayaux, P., & Lambin, E. F. (1995). Estimation of tropical forest area from coarse spatial resolution data: A two-step correction function for proportional errors due to spatial aggregation. *Remote Sensing of Environment*, 53, 1–15.
- Pareto, V. (1906). *Manuale d'Economia Politica*. Milano: Società Editrice Libreria.
- Plourde, L., & Congalton, R. (2003). Sampling method and sample placement: How do they affect the accuracy of remotely sensed maps? *Photogrammetric Engineering and Remote Sensing*, 69, 289–297.
- Rommel, T. K., & Perera, A. H. (2002). Accuracy of discontinuous binary surfaces: A case study using boreal forest fires. *International Journal of Geographical Information Science*, 16, 287–298.
- Roy, D., Lewis, P. E., & Justice, C. O. (2002). Burned area mapping using multi-temporal moderate spatial resolution data—a bi-directional reflectance model-based expectation approach. *Remote Sensing of the Environment*, 83, 263–286.
- Simon, M., Plummer, S., Fierens, F., Hoeltzemann, J. J., & Arino, O. (2004). Burnt area detection at global scale using ATSR-2: The Globscar products and their qualification. *Journal of Geophysical Research* (In press).
- Smith, J. H., Stehman, S. V., Wickham, J. D., & Yang, W. (2003). Effects of landscape characteristics on land-cover class accuracy. *Remote Sensing of Environment*, 84, 342–349.
- Smits, P. C., Dellepiane, S. G., & Showengerdt, R. A. (1999). Quality assessment of image classification algorithms for land-cover mapping: A review and a proposal for a cost-based approach. *International Journal of Remote Sensing*, 20(8), 1461–1486.
- Stehman, S. V. (1997). Selecting and interpreting measures of thematic classification accuracy. *Remote Sensing of Environment*, 62, 77–89.
- Story, M., & Congalton, R. G. (1986). Accuracy assessment: A user's perspective. *Photogrammetric Engineering and Remote Sensing*, 52, 397–399.
- Tansey, K., Grégoire, J. -M., Stroppiana, D., Sousa, A., Silva, J. M. N., Pereira, J. M. C., Boschetti, L., Maggi, M., Brivio, P. A., Fraser, R., Flasse, S., Ershov, D., Binaghi, E., Graetz, D., & Peduzzi, P. (2004). Vegetation burning in the year 2000: Global burned area estimates from SPOT VEGETATION data. *Journal of Geophysical Research* (In press).
- Thomlinson, J. R., Bolstad, P. V., & Cohen, W. B. (1999). Coordinating methodologies for scaling landcover classifications from site-specific to global: Steps toward validation global map products. *Remote Sensing of Environment*, 70, 16–28.
- Woodcock, C., & Strahler, A. (1987). The factor of scale in remote sensing. *Remote Sensing of Environment*, 20, 311–332.